

Machine Learning 1: Linear Regression

Stefano Ermon













































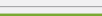




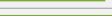




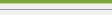
March 31, 2016

Plan for today

Plan for today:

- Supervised Machine Learning: linear regression

Renewable electricity generation in the U.S

	Hydropower	Solar ¹	Wind	Geothermal	Biomass	Total Renewables
2004	 6.7%	0.0%	 0.4%	 0.4%	 1.3%	 8.8%
2005	 6.7%	0.0%	 0.4%	 0.4%	 1.3%	 8.8%
2006	 7.1%	0.0%	 0.7%	 0.4%	 1.3%	 9.5%
2007	 5.9%	0.0%	 0.8%	 0.4%	 1.3%	 8.5%
2008	 6.2%	0.1%	 1.3%	 0.4%	 1.3%	 9.3%
2009	 6.9%	0.1%	 1.9%	 0.4%	 1.4%	 10.6%
2010	 6.3%	0.1%	 2.3%	 0.4%	 1.4%	 10.4%
2011	 7.8%	0.2%	 2.9%	 0.4%	 1.4%	 12.6%
2012	 6.8%	0.3%	 3.4%	 0.4%	 1.4%	 12.4%
2013	 6.6%	0.5%	 4.1%	 0.4%	 1.5%	 13.1%
2014	 6.3%	0.8%	 4.4%	 0.4%	 1.6%	 13.5%

Source: Renewable energy data book, NREL


Challenges for the grid

- Wind and solar are intermittent
- We will need traditional power plants when the wind stops
 - Many power plants (e.g., nuclear) cannot be easily turned on/off or quickly ramped up/down
- With more accurate forecasts, wind and solar power become more efficient alternatives
 - A few years ago, Xcel Energy (Colorado) ran ads opposing a proposal that it use 10% of renewable sources
 - Thanks to wind forecasting (ML) algorithms developed at NCAR, they now aim for 30 percent. Accurate forecasting saved the utility \$6-\$10 million per year

Motivation

- Solar and wind are intermittent
- Can we accurately forecast how much energy will we consume tomorrow?
 - Difficult to estimate from “a priori” models
 - But, we have lots of data from which to build a model


Typical electricity consumption



Data: PJM <http://www.pjm.com>

Predict peak demand from high temperature

- What will peak demand be tomorrow?
- If we know something else about tomorrow (like the high temperature), we can use this to *predict* peak demand




Data: PJM, Weather Underground (summer months, June-August)

A simple model

- A linear model that predicts demand:

$$\text{predicted peak demand} = \theta_1 \cdot (\text{high temperature}) + \theta_2$$



- *Parameters of model:* $\theta_1, \theta_2 \in \mathbb{R}$ ($\theta_1 = 0.046$, $\theta_2 = -1.46$)

A simple model

- We can use a model like this to make predictions
- What will be the peak demand tomorrow?
 - I know from weather report that high temperature will be 80°F (ignore, for the moment, that this too is a prediction)
- Then predicted peak demand is:

$$\theta_1 \cdot 80 + \theta_2 = 0.046 \cdot 80 - 1.46 = 2.19 \text{ GW}$$

Formal problem setting

- **Input:** $x_i \in \mathbb{R}^n$, $i = 1, \dots, m$
 - E.g.: $x_i \in \mathbb{R}^1 = \{\text{high temperature for day } i\}$
- **Output:** $y_i \in \mathbb{R}$ (*regression* task)
 - E.g.: $y_i \in \mathbb{R} = \{\text{peak demand for day } i\}$
- **Model Parameters:** $\theta \in \mathbb{R}^k$
- **Predicted Output:** $\hat{y}_i \in \mathbb{R}$

$$\text{E.g.: } \hat{y}_i = \theta_1 \cdot x_i + \theta_2$$

- For convenience, we define a function that maps inputs to *feature vectors*

$$\phi : \mathbb{R}^n \rightarrow \mathbb{R}^k$$

- For example, in our task above, if we define

$$\phi(x_i) = \begin{bmatrix} x_i \\ 1 \end{bmatrix} \quad (\text{here } n = 1, k = 2)$$

then we can write

$$\hat{y}_i = \sum_{j=1}^k \theta_j \cdot \phi_j(x_i) \equiv \theta^T \phi(x_i)$$

Loss functions

- Want a model that performs “well” on the data we have

$$\text{i.e., } \hat{y}_i \approx y_i, \quad \forall i$$

- We measure “closeness” of \hat{y}_i and y_i using *loss function*

$$\ell : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}_+$$

- Example: squared loss

$$\ell(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

Finding model parameters, and optimization

- Want to find model parameters such that minimize sum of costs over all input/output pairs

$$J(\theta) = \sum_{i=1}^m \ell(\hat{y}_i, y_i) = \sum_{i=1}^m (\theta^T \phi(x_i) - y_i)^2$$

- Write our objective formally as

$$\underset{\theta}{\text{minimize}} \quad J(\theta)$$

simple example of an *optimization problem*; these will dominate our development of algorithms throughout the course

How do we optimize a function

- Search algorithm: Start with an initial guess for θ . Keep changing θ (by a little bit) to reduce $J(\theta)$
- Animation <https://www.youtube.com/watch?v=vWFjqgb-y1Q>

Gradient descent

- Search algorithm: Start with an initial guess for θ . Keep changing θ (by a little bit) to reduce $J(\theta)$

$$J(\theta) = \sum_{i=1}^m \ell(\hat{y}_i, y_i) = \sum_{i=1}^m (\theta^T \phi(x_i) - y_i)^2$$

- Gradient descent: $\theta_j = \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}$, for all j

$$\begin{aligned} \frac{\partial J}{\partial \theta_j} &= \frac{\partial \sum_{i=1}^m (\theta^T \phi(x_i) - y_i)^2}{\partial \theta_j} = \sum_{i=1}^m \frac{\partial (\theta^T \phi(x_i) - y_i)^2}{\partial \theta_j} \\ &= \sum_{i=1}^m 2(\theta^T \phi(x_i) - y_i) \frac{\partial (\theta^T \phi(x_i) - y_i)}{\partial \theta_j} \\ &= \sum_{i=1}^m 2(\theta^T \phi(x_i) - y_i) \phi(x_i)_j \end{aligned}$$

Gradient descent

- Repeat until “convergence”:

$$\theta_j = \theta_j - \alpha \sum_{i=1}^m 2(\theta^T \phi(x_i) - y_i) \phi(x_i)_j, \text{ for all } j$$

- Demo:
<https://lukaszkujava.github.io/gradient-descent.html>
- Stochastic gradient descent

- Let's write $J(\theta)$ a little more compactly using matrix notation; define

$$\Phi \in \mathbb{R}^{m \times k} = \begin{bmatrix} - & \phi(x_1)^T & - \\ - & \phi(x_2)^T & - \\ & \vdots & \\ - & \phi(x_m)^T & - \end{bmatrix}, \quad y \in \mathbb{R}^m = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$$


then

$$J(\theta) = \sum_{i=1}^m (\theta^T \phi(x_i) - y_i)^2 = \|\Phi\theta - y\|_2^2$$


($\|z\|_2$ is ℓ_2 norm of a vector: $\|z\|_2 \equiv \sqrt{\sum_{i=1}^m z_i^2} = \sqrt{z^T z}$)

- Called *least-squares* objective function

- How do we optimize a function? 1-D case ($\theta \in \mathbb{R}$):



$$J(\theta) = \theta^2 - 2\theta - 1$$



$$\frac{dJ}{d\theta} = 2\theta - 2$$

$$\theta^* \text{ minimum} \implies \left. \frac{dJ}{d\theta} \right|_{\theta^*} = 0$$

$$\implies 2\theta^* - 2 = 0$$

$$\implies \theta^* = 1$$

- Multi-variate case: $\theta \in \mathbb{R}^k$, $J : \mathbb{R}^k \rightarrow \mathbb{R}$

Generalized condition: $\nabla_{\theta} J(\theta)|_{\theta^*} = 0$

- $\nabla_{\theta} J(\theta)$ denotes *gradient* of J with respect to θ

$$\nabla_{\theta} J(\theta) \in \mathbb{R}^k \equiv \begin{bmatrix} \frac{\partial J}{\partial \theta_1} \\ \frac{\partial J}{\partial \theta_2} \\ \vdots \\ \frac{\partial J}{\partial \theta_k} \end{bmatrix}$$

- Some important rules and common gradient

$$\nabla_{\theta}(af(\theta) + bg(\theta)) = a\nabla_{\theta}f(\theta) + b\nabla_{\theta}g(\theta), \quad (a, b \in \mathbb{R})$$

$$\nabla_{\theta}(\theta^T A \theta) = (A + A^T)\theta, \quad (A \in \mathbb{R}^{k \times k})$$

$$\nabla_{\theta}(b^T \theta) = b, \quad (b \in \mathbb{R}^k)$$

- Optimizing least-squares objective

$$\begin{aligned} J(\theta) &= \|\Phi\theta - y\|_2^2 \\ &= (\Phi\theta - y)^T (\Phi\theta - y) \\ &= \theta^T \Phi^T \Phi \theta - 2y^T \Phi \theta + y^T y \end{aligned}$$

- Using the previous gradient rules

$$\begin{aligned} \nabla_{\theta} J(\theta) &= \nabla_{\theta} (\theta^T \Phi^T \Phi \theta - 2y^T \Phi \theta + y^T y) \\ &= \nabla_{\theta} (\theta^T \Phi^T \Phi \theta) - 2\nabla_{\theta} (y^T \Phi \theta) + \nabla_{\theta} (y^T y) \\ &= 2\Phi^T \Phi \theta - 2\Phi^T y \end{aligned}$$

- Setting gradient equal to zero

$$2\Phi^T \Phi \theta^* - 2\Phi^T y = 0 \iff \theta^* = (\Phi^T \Phi)^{-1} \Phi^T y$$

known as the *normal equations*

- Let's see how this looks in MATLAB code

```
X = load(high_temperature.txt);  
y = load(peak_demand.txt);  
n = size(X,2);  
m = size(X,1);  
Phi = [X ones(m,1)];  
theta = inv(Phi' * Phi) * Phi' * y;  
  
theta =  
    0.0466  
   -1.4600
```

- The normal equations are so common that MATLAB has a special operation for them

```
% same as inv(Phi' * Phi) * Phi' * y  
theta = Phi \ y;
```

Higher-dimensional inputs

- Input: $x \in \mathbb{R}^2 = \begin{bmatrix} \text{temperature} \\ \text{hour of day} \end{bmatrix}$
- Output: $y \in \mathbb{R} = \text{demand}$

- Features: $\phi(x) \in \mathbb{R}^3 = \begin{bmatrix} \text{temperature} \\ \text{hour of day} \\ 1 \end{bmatrix}$

- Same matrices as before

$$\Phi \in \mathbb{R}^{m \times k} = \begin{bmatrix} - & \phi(x_1)^T & - \\ & \vdots & \\ - & \phi(x_m)^T & - \end{bmatrix}, \quad y \in \mathbb{R}^m = \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix}$$

- Same solution as before

$$\theta \in \mathbb{R}^3 = (\Phi^T \Phi)^{-1} \Phi^T y$$

